

# Lecture 2:

## Pixels and Filters

# What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

# What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

# Types of Images

## Binary



# Types of Images

**Binary**



**Gray Scale**



# Types of Images

**Binary**




**Gray Scale**



**Color**




# Binary image representation



Slide credit: Ulas Bagci


# Grayscale image representation



Slide credit: Ulas Bagci



# Color Image - one channel



Slide credit: Ulas Bagci



# Color image representation



Slide credit: Ulas Bagci

# Images are sampled

What happens when we zoom into the images we capture?



# Resolution


is a **sampling** parameter, defined in dots per inch (DPI) or equivalent measures of spatial pixel density, and its standard value for recent screen technologies is 72 dpi



Slide credit: Ulas Bagci

# Images are Sampled and Quantized

- An image contains discrete number of pixels
  - A simple example
  - Pixel value:
    - “grayscale”  
(or “intensity”):  $[0, 255]$





# Images are Sampled and Quantized

- An image contains discrete number of pixels
    - A simple example
    - Pixel value:
      - “grayscale”  
(or “intensity”):  $[0, 255]$
      - “color”
        - RGB:  $[R, G, B]$
        - Lab:  $[L, a, b]$
        - HSV:  $[H, S, V]$
- 
- The image is a pixelated representation of a face. Three arrows point to specific pixels with their corresponding RGB values:
- Top arrow:  $[90, 0, 53]$  (purple)
  - Bottom-left arrow:  $[213, 60, 67]$  (red)
  - Right arrow:  $[249, 215, 203]$  (red)

# What we will learn today?


- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

# Histogram

- Histogram captures the distribution of gray levels in the image.
- How frequently each gray level occurs in the image





# What we will learn today?

- Image sampling and quantization
- Image histograms
- **Images as functions**
- Linear systems (filters)
- Convolution and correlation


Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

# Images as discrete functions

- Images are usually **digital (discrete)**:
  - **Sample** the 2D space on a regular grid
- Represented as a matrix of integer values

pixel



$j$	62	79	23	119	120	05	4	0
$i$	10	10	9	62	12	78	34	0
	10	58	197	46	46	0	0	48
	176	135	5	188	191	68	0	49
	2	1	1	29	26	37	0	77
	0	89	144	147	187	102	62	208
	255	252	0	166	123	62	0	31
	166	63	127	17	1	0	99	30

# Images as coordinates

## Cartesian coordinates

$$I[n, m] = \begin{bmatrix} \ddots & & \vdots & & \\ & I[-1, 1] & I[0, 1] & I[1, 1] & \\ \dots & I[-1, 0] & \underline{I[0, 0]} & I[1, 0] & \dots \\ & I[-1, -1] & I[0, -1] & I[1, -1] & \\ & & \vdots & & \ddots \end{bmatrix}$$

$I[n, m]$  is the notation for discrete functions


# Images as functions

- **An Image** as a function  $f$  from  $\mathbb{R}^2$  to  $\mathbb{R}^M$ :
  - $I(x, y)$  gives the **intensity** at position  $(x, y)$
  - Defined over a rectangle, with a finite range:

$$I: [a,b] \times [c,d] \rightarrow [0,255]$$

Domain  
support

range



# Images as functions

- **An Image** as a function  $f$  from  $\mathbb{R}^2$  to  $\mathbb{R}^M$ :
  - $I(x, y)$  gives the **intensity** at position  $(x, y)$
  - Defined over a rectangle, with a finite range:

$$I: \underbrace{[a,b] \times [c,d]}_{\substack{\text{Domain} \\ \text{support}}} \rightarrow \underbrace{[0,255]}_{\text{range}}$$

- A color image:  $I(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$

# What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Image filtering

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

# Systems and Filters

## Filtering:

- Forming a new image whose pixel values are transformed from original pixel values

## Goals:

- Goal is to extract useful information from images, or transform images into another domain where we can modify/enhance image properties
  - Features (edges, corners, blobs...)
  - super-resolution; in-painting; de-noising


## De-noising



Salt and pepper noise



## Super-resolution



## In-painting



Bertamio et al



# Image filtering

- Image filters in spatial domain
  - Filter is a mathematical operation of a grid of numbers
  - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
  - Filtering is a way to modify the frequencies of images
  - Denoising, sampling, image compression

# Image filtering

- Image filtering:
  - Compute function of local neighborhood at each position

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

# Image filtering

- Image filtering:
  - Compute function of local neighborhood at each position

`h=output`                      `f=filter`      `I=image`

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

2d coords=`k,l`      2d coords=`m,n`

[ ]

[ ]

[ ]

# Example: box filter

$$\frac{1}{9} f[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$


$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10							

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20						

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$



# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				
						?			
				50					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering

$$f[\cdot, \cdot] \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$\frac{1}{9} f[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

# Box Filter


What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)
- Why does it sum to one?

$$\frac{1}{9} f[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

# Smoothing with box filter



# Image filtering

- Image filtering:
  - Compute function of local neighborhood at each position

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

- Really important!
  - Enhance images
    - Denoise, resize, increase contrast, etc.
  - Extract information from images
    - Texture, edges, distinctive points, etc.
  - Detect patterns
    - Template matching



# Think-Pair-Share time



1.

0	0	0
0	1	0
0	0	0

2.

0	0	0
0	0	1
0	0	0

3.

1	0	-1
2	0	-2
1	0	-1

4.

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

# 1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

# 1. Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

## 2. Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

## 2. Practice with linear filters




Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

### 3. Practice with linear filters




1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge  
(absolute value)

### 3. Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)

# 4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)



# 4. Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

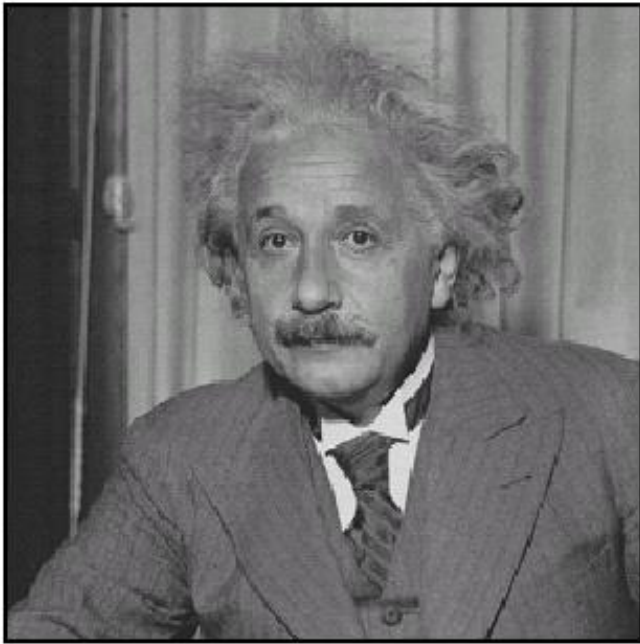
1	1	1
1	1	1
1	1	1



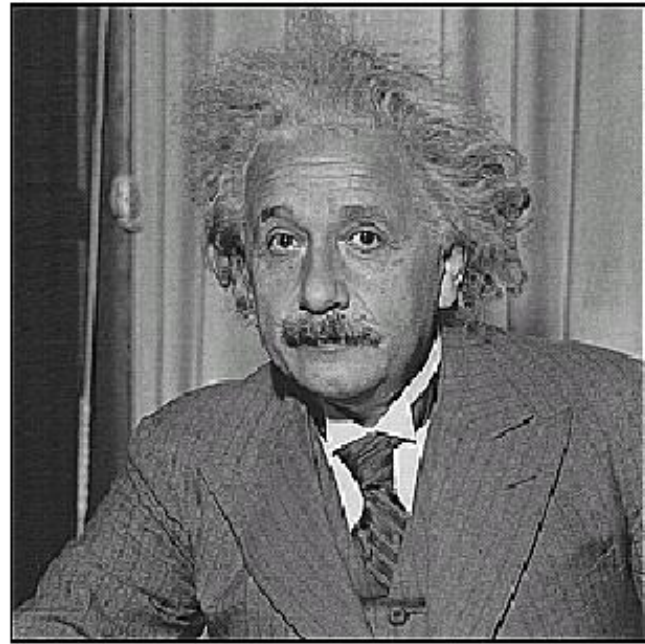
## Sharpening filter

- Accentuates differences with local average

## 4. Practice with linear filters



before



after

# Correlation and Convolution

- 2d correlation

$$h[m,n] = \sum_{k,l} f[k,l] I[m+k,n+l]$$

`h=filter2(f,I);` or `h=imfilter(I,f);`

- 2d convolution

$$h[m,n] = \sum_{k,l} f[k,l] I[m-k,n-l]$$



`h=conv2(f,I);` or `h=imfilter(I,f,'conv');`

Correlation and convolution are identical when the filter is symmetric.

# 1D Discrete convolution (symbol: $*$ )

We are going to convolve a function **f** with a filter **h**.

$$g[n] = \sum_k f[k]h[n - k]$$






# 1D Discrete convolution (symbol: $*$ )

We are going to convolve a function **f** with a filter **h**.


$$g[n] = \sum_k f[k]h[n - k]$$

We first need to calculate  $h[n-k, m-l]$






# Discrete convolution (symbol: $*$ )

We are going to convolve a function **f** with a filter **h**.






# Discrete convolution (symbol: $*$ )

We are going to convolve a function **f** with a filter **h**.



# Discrete convolution (symbol: $*$ )




We are going to convolve a function  $\mathbf{f}$  with a filter  $\mathbf{h}$ .






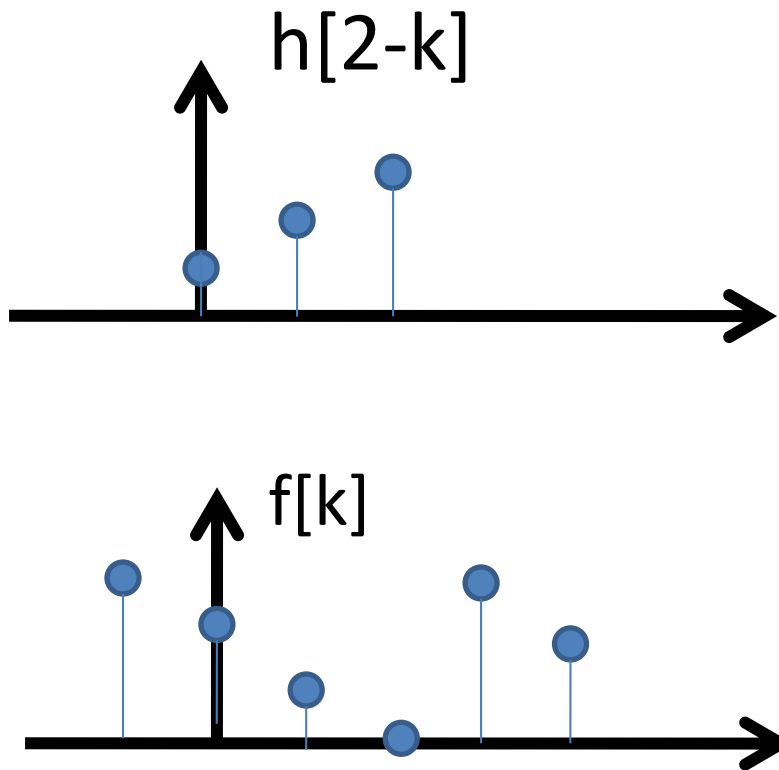
# Discrete convolution (symbol: $*$ )

We are going to convolve a function  $\mathbf{f}$  with a filter  $\mathbf{h}$ .




# Discrete convolution (symbol: $*$ )

We are going to convolve a function  $\mathbf{f}$  with a filter  $\mathbf{h}$ .



# Discrete convolution (symbol: $*$ )

We are going to convolve a function **f** with a filter **h**.




# 2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter( $h[,]$ ) that is 3x3. and an image ( $f[,]$ ) that is 7x7.


n

# 2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter(h[,]) that is 3x3. and an image (f[,]) that is 7x7.


n

# 2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter( $h[,]$ ) that is 3x3. and an image ( $f[,]$ ) that is 7x7.


n

# 2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter( $h[,]$ ) that is 3x3. and an image ( $f[,]$ ) that is 7x7.


n

# 2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter( $h[,]$ ) that is 3x3. and an image ( $f[,]$ ) that is 7x7.

n




# 2D convolution

2D convolution is very similar to 1D.


- The main difference is that we now have to iterate over 2 axis instead of 1.

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter(h[,]) that is 3x3. and an image (f[,]) that is 7x7.

# 2D convolution example



# 2D convolution example

1	2	1	
0	0	0	3
-1	-2	-1	6
	7	8	9

$$\begin{aligned}
 &= x[-1,-1] \cdot h[1,1] + x[0,-1] \cdot h[0,1] + x[1,-1] \cdot h[-1,1] \\
 &\quad + x[-1,0] \cdot h[1,0] + x[0,0] \cdot h[0,0] + x[1,0] \cdot h[-1,0] \\
 &\quad + x[-1,1] \cdot h[1,-1] + x[0,1] \cdot h[0,-1] + x[1,1] \cdot h[-1,-1] \\
 &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 0 \cdot 0 + 1 \cdot 0 + 2 \cdot 0 + 0 \cdot (-1) + 4 \cdot (-2) + 5 \cdot (-1) = -13
 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

# 2D convolution example

1	2	1
0	0	0
-1	-2	-1
7	8	9

$$\begin{aligned}
 &= x[0,-1] \cdot h[1,1] + x[1,-1] \cdot h[0,1] + x[2,-1] \cdot h[-1,1] \\
 &\quad + x[0,0] \cdot h[1,0] + x[1,0] \cdot h[0,0] + x[2,0] \cdot h[-1,0] \\
 &\quad + x[0,1] \cdot h[1,-1] + x[1,1] \cdot h[0,-1] + x[2,1] \cdot h[-1,-1] \\
 &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 0 + 2 \cdot 0 + 3 \cdot 0 + 4 \cdot (-1) + 5 \cdot (-2) + 6 \cdot (-1) = -20
 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

# 2D convolution example

		1	2	1
1	0	2	0	0
4	-1	5	-2	-1
7	8	9		

$$\begin{aligned}
 &= x[1,-1] \cdot h[1,1] + x[2,-1] \cdot h[0,1] + x[3,-1] \cdot h[-1,1] \\
 &\quad + x[1,0] \cdot h[1,0] + x[2,0] \cdot h[0,0] + x[3,0] \cdot h[-1,0] \\
 &\quad + x[1,1] \cdot h[1,-1] + x[2,1] \cdot h[0,-1] + x[3,1] \cdot h[-1,-1] \\
 &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 2 \cdot 0 + 3 \cdot 0 + 0 \cdot 0 + 5 \cdot (-1) + 6 \cdot (-2) + 0 \cdot (-1) = -17
 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

# 2D convolution example

1	2	1	
	1	2	3
0	0	0	
	4	5	6
-1	-2	-1	
	7	8	9

$$\begin{aligned}
 &= x[-1,0] \cdot h[1,1] + x[0,0] \cdot h[0,1] + x[1,0] \cdot h[-1,1] \\
 &\quad + x[-1,1] \cdot h[1,0] + x[0,1] \cdot h[0,0] + x[1,1] \cdot h[-1,0] \\
 &\quad + x[-1,2] \cdot h[1,-1] + x[0,2] \cdot h[0,-1] + x[1,2] \cdot h[-1,-1] \\
 &= 0 \cdot 1 + 1 \cdot 2 + 2 \cdot 1 + 0 \cdot 0 + 4 \cdot 0 + 5 \cdot 0 + 0 \cdot (-1) + 7 \cdot (-2) + 8 \cdot (-1) = -18
 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

# 2D convolution example

1 1	2 2	1 3
0 4	0 5	0 6
-1 7	-2 8	-1 9

$$\begin{aligned} &= x[0,0] \cdot h[1,1] + x[1,0] \cdot h[0,1] + x[2,0] \cdot h[-1,1] \\ &\quad + x[0,1] \cdot h[1,0] + x[1,1] \cdot h[0,0] + x[2,1] \cdot h[-1,0] \\ &\quad + x[0,2] \cdot h[1,-1] + x[1,2] \cdot h[0,-1] + x[2,2] \cdot h[-1,-1] \\ &= 1 \cdot 1 + 2 \cdot 2 + 3 \cdot 1 + 4 \cdot 0 + 5 \cdot 0 + 6 \cdot 0 + 7 \cdot (-1) + 8 \cdot (-2) + 9 \cdot (-1) = -24 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

# 2D convolution example

1	1	2	1
4	0	0	0
7	-1	-2	-1

$$\begin{aligned}
 &= x[1,0] \cdot h[1,1] + x[2,0] \cdot h[0,1] + x[3,0] \cdot h[-1,1] \\
 &\quad + x[1,1] \cdot h[1,0] + x[2,1] \cdot h[0,0] + x[3,1] \cdot h[-1,0] \\
 &\quad + x[1,2] \cdot h[1,-1] + x[2,2] \cdot h[0,-1] + x[3,2] \cdot h[-1,-1] \\
 &= 2 \cdot 1 + 3 \cdot 2 + 0 \cdot 1 + 5 \cdot 0 + 6 \cdot 0 + 0 \cdot 0 + 8 \cdot (-1) + 9 \cdot (-2) + 0 \cdot (-1) = -18
 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output



# Convolution in 2D - examples



Original



•0	•0	•0
•0	•1	•0
•0	•0	•0



# Convolution in 2D - examples



Original



•0	•0	•0
•0	•1	•0
•0	•0	•0



Filtered  
(no change)

# Convolution in 2D - examples



Original



•0	•0	•0
•0	•0	•1
•0	•0	•0



# Convolution in 2D - examples



Original



•0	•0	•0
•0	•0	•1
•0	•0	•0



Shifted right  
By 1 pixel

# Convolution in 2D - examples



Original

\*

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

=

?

# Convolution in 2D - examples



Original



$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1



Blur (with a  
box filter)

# Convolution in 2D - examples



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

—

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

= ?

(Note that filter sums to 1)

“details of the image”

•0	•0	•0
•0	•1	•0
•0	•0	•0

+

•0	•0	•0
•0	•1	•0
•0	•0	•0

—


$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1


- What does blurring take away?



—




=



- Let's add it back:



+ a



=





# Convolution in 2D – Sharpening filter



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

–

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

=



**Sharpening filter:** Accentuates differences with local average

# Convolution properties

Commutative:  $a * b = b * a$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality, e.g., image edges

Associative:  $a * (b * c) = (a * b) * c$

- Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$

# Convolution properties

## Proof of associativity of convolution

$$\begin{aligned}((f \star g) \star h)(n) &= \sum_{k=0}^n (f \star g)(k) h(n-k) \\&= \sum_{k=0}^n \sum_{l=0}^k f(l) g(k-l) h(n-k) \\&= \sum_{l=0}^n \sum_{k=l}^n f(l) g(k-l) h(n-k) \\&= \sum_{l=0}^n \sum_{k=0}^{n-l} f(l) g(k) h(n-k-l) \\&= \sum_{l=0}^n f(l) (g \star h)(n-l) \\&= (f \star (g \star h))(n)\end{aligned}$$

# Convolution properties

Commutative:  $a * b = b * a$

- Conceptually no difference between filter and signal
- But particular filtering implementations might break this equality, e.g., image edges

Associative:  $a * (b * c) = (a * b) * c$


- Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
- Correlation is not associative
- Why important?

# Convolution properties

- Commutative:  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
  - But particular filtering implementations might break this equality, e.g., image edges
- Associative:  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
  - Correlation is not associative (rotation effect)
  - Why important?
- Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out:  $ka * b = a * kb = k(a * b)$


# Important filter: Gaussian

- Weight contributions of neighboring pixels by nearness




x				
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5,  $\sigma = 1$




$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Smoothing with Gaussian filter



# Smoothing with box filter





# Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
  - Images become more smooth
- Gaussian convolved with Gaussian...
  - ...is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolution two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

# Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{x^2}{2\sigma^2} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{y^2}{2\sigma^2} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Separability example

2D convolution  
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution  
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution  
along the remaining column:

# Separability

Why is separability useful in practice?

# Separability

Why is separability useful in practice?

MxN image, PxQ filter

- 2D convolution:  $\sim MN PQ$  multiply-adds
- Separable 2D:  $\sim MN(P+Q)$  multiply-adds

Speed up =  $PQ/(P+Q)$

9x9 filter =  $\sim 4.5x$  faster


# Practical matters

## How big should the filter be?

- Values at edges should be near zero
- Gaussians have infinite extent...
- Rule of thumb for Gaussian: set filter half-width to about  $3 \sigma$

# Practical matters

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



# Convolution in Convolutional Neural Networks

- Convolution is the basic operation in CNNs
- Learning convolution kernels allows us to learn which `features' provide useful information in images.



# **NON-LINEAR FILTERS**

# Median filters

- Operates over a window by selecting the median intensity in the window.
- ‘Rank’ filter as based on ordering of gray levels
  - E.G., min, max, range filters

# Image filtering - mean

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Image filtering - mean

$$f[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$I[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[\cdot, \cdot]$

	0	10	20	30	30				
				50					

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

# Median filter?

 $I[.,.]$ 

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $h[.,.]$ 

				?					

# Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?

# Noisy Jack – Salt and Pepper



# Mean Jack – 3 x 3 filter





Very Mean Jack – 11 x 11 filter



# Noisy Jack – Salt and Pepper



# Median Jack – 3 x 3



# Very Median Jack – 11 x 11



# Median filters

- Operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?

# Review: questions

1. Write down a 3x3 filter that both:
  - Returns a positive value if the average value of the 4-adjacent neighbors is less than the center,
  - Returns a negative value otherwise.

# Review: questions

1. Write down a 3x3 filter that both:

- Returns a positive value if the average value of the 4-adjacent neighbors is less than the center,
  - Returns a negative value otherwise.
- $$\begin{bmatrix} 0 & -1/4 & 0; \\ -1/4 & 1 & -1/4; \\ 0 & -1/4 & 0 \end{bmatrix}$$